


Model Cards Schema

LinkML schema · multi-method generation · D4D harmonization · hybrid + LLM rubric evaluation



Portfolio: climate-extended · climate-forecasting · sentiment · densenet121 · subcell-saprot · minimal fixtures

Zoom · Model Card Generation

Model docs (HF Hub · GitHub · papers · DOIs) → Model Card YAML

Sources

HuggingFace Hub README · GitHub MODEL_CARD.md · papers (arXiv) · DOIs (ARTL MCP) · KOGUT docs → preprocessed + concatenated per project

Generation methods

★ mc-agent

parallel sub-agents per section
deterministic · highest quality
/mc-agent skill

mc-assistant

interactive, human-in-the-loop
good for refinement
/mc-assistant skill

GitHub @modelcardassistant

triggered by issue / PR mention
URL / DOI ingest modes
opens PR with card + portfolio diff

HF Hub direct

README → Model Card
densenet121, SaProt, ...

Canonical / curated

DOE climate-extended
gold reference

Floor anchors (fixtures)

minimal · minimal-viable
rubric calibration

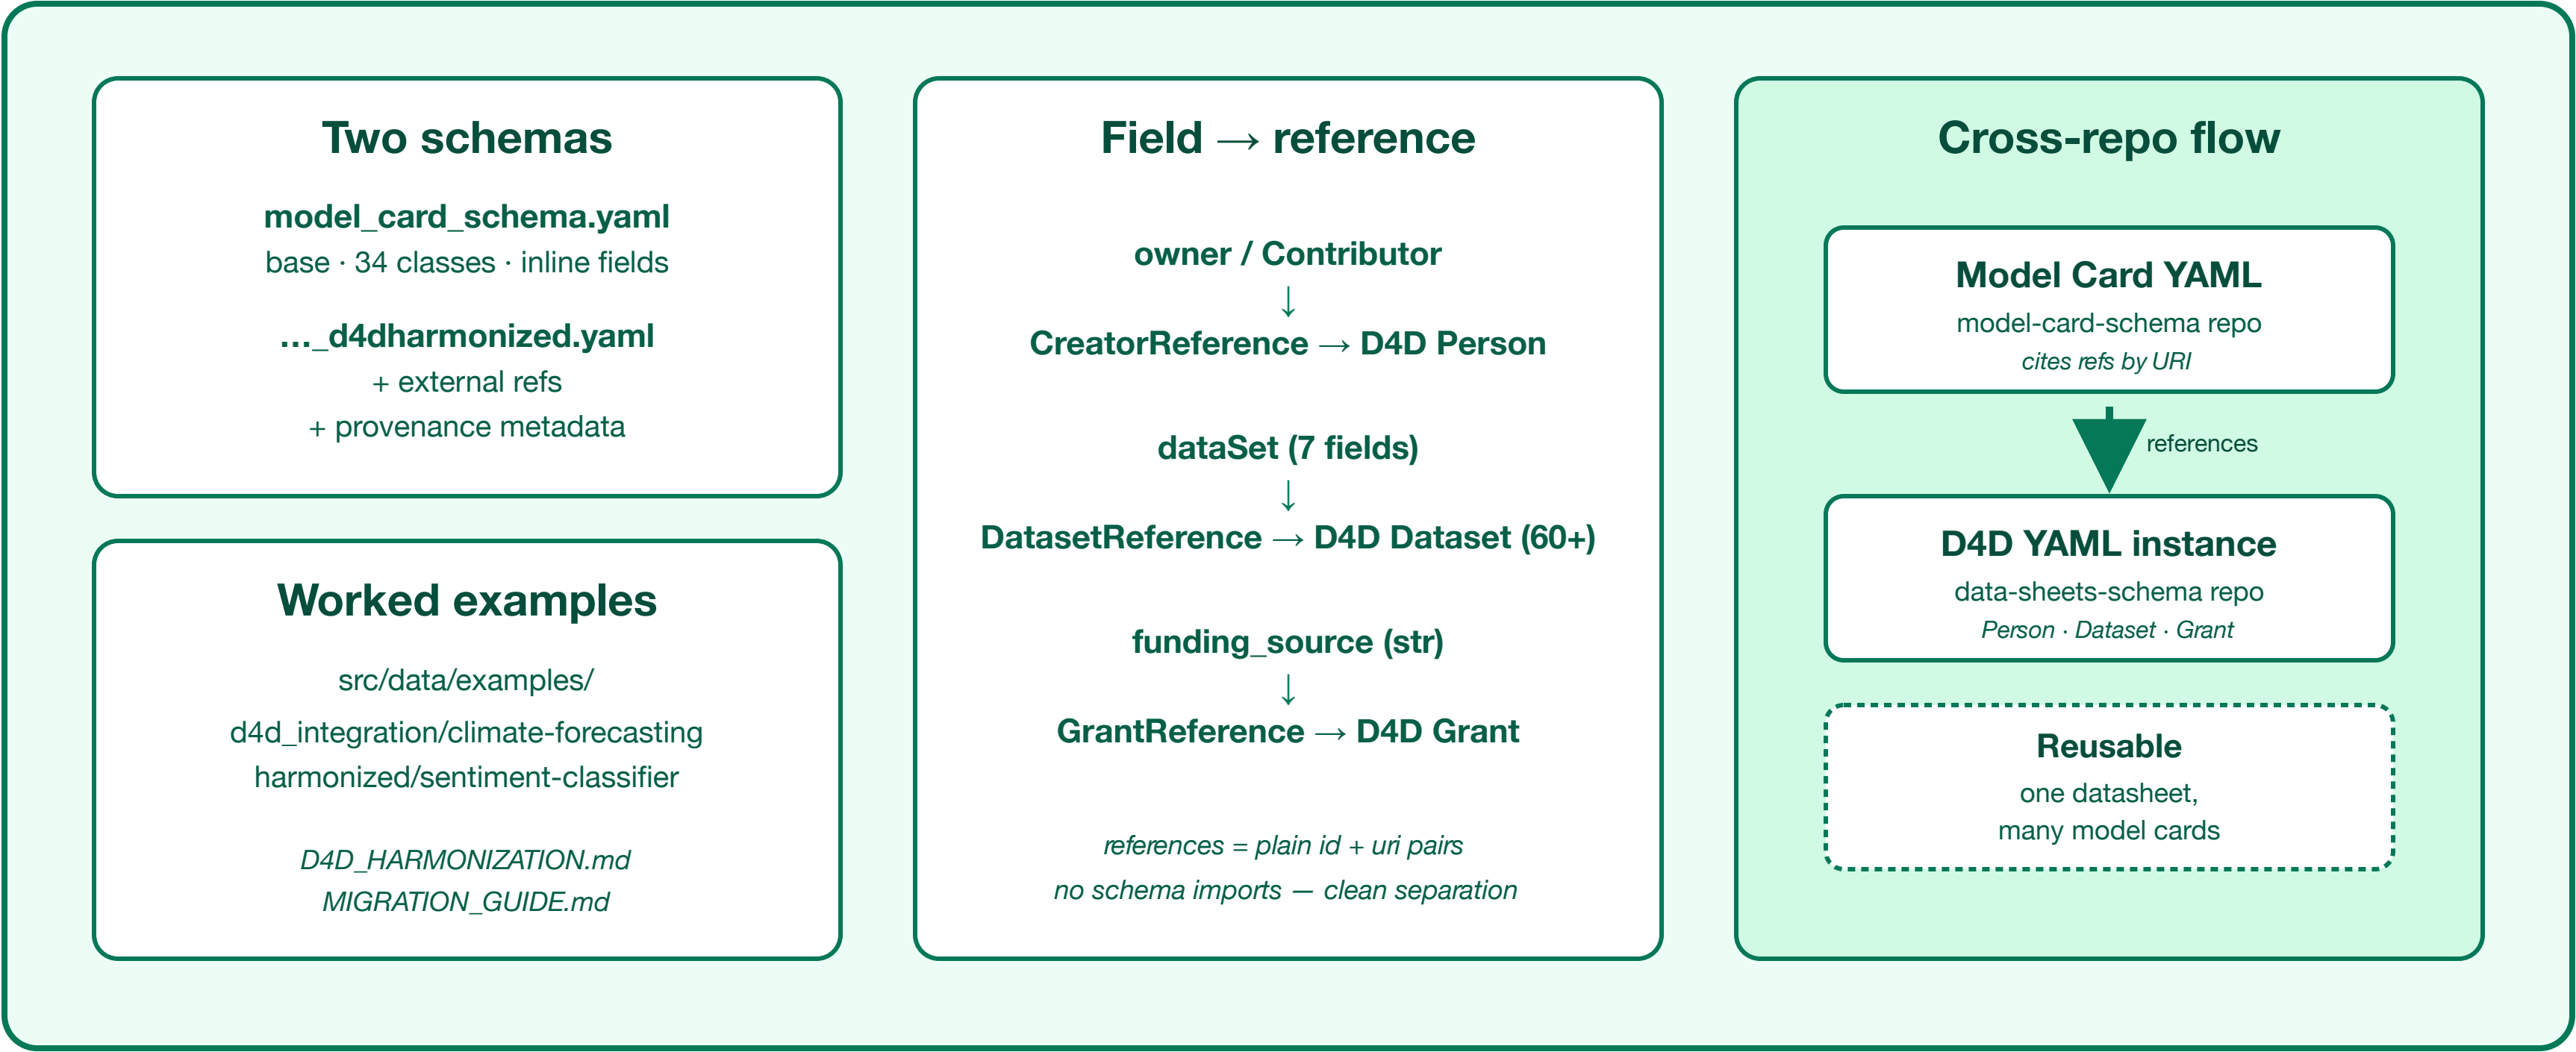
make targets

gen-project · test · lint
gendoc · serve
regenerates JSON / OWL / Python / ...

Model Card YAML · data/model_cards_assistant/, ..._concatenated/, ..._individual/ · blocked by CI gate before merge

Zoom · D4D Harmonization & External References

Lightweight inline fields → ID + URI references to comprehensive D4D instances



Replaces shallow inline metadata with FAIR, governance-grade dataset / creator / grant documentation

Zoom · Rubric Evaluation (Hybrid + LLM)

Two axes: which rubric (10 vs 20) × which evaluator (hybrid rule-based vs LLM judge)

rubric10

10 elements × 5 sub
max 50 pts · coarse

rubric20

20 questions, 4 cats:
Struct · Meta · Tech · FAIR

Hybrid (rule-based)

deterministic · free · <1 s · runs in CI gate

r10 × hybrid

`batch_evaluate_mc_rubric10_hybrid.py`

fast sanity check

runs on every PR

r20 × hybrid

`batch_evaluate_mc_rubric20_hybrid.py`

bakes in semantic findings

Q18 leakage · Q19 bias ↔ tradeoffs

LLM judge

`mc-rubric{10,20}(-semantic)` · Sonnet 4.5 · T=0.0

r10 × LLM

`mc-rubric10 / mc-rubric10-semantic`

deeper quality signal

on-demand · per card cost

r20 × LLM

`mc-rubric20 / mc-rubric20-semantic`

gold-standard FAIR audit

surfaces rules to encode into hybrid

CI & portfolio

`mc-rubric20-gate.yml` · blocking PR check, hybrid only (deterministic, no API key)

`mc-rubric20-portfolio-audit.yml` · weekly cron · rebuilds `portfolio.html` · badges · `common_issues.md` · opens PR if anything changed

Outputs: `data/evaluation/{badges, all, extended, harmonized, hf_hub}/` · per-card SVG badge (rubric × evaluator) in README

Model Cards: Generation, D4D Harmonization, and Evaluation

LinkML schema · multi-method extraction · external D4D references · hybrid + LLM rubric scoring · CI-gated portfolio

